

workbook (2)

January 9, 2023

1 Biomedical Image Analysis

```
[4]: import os
import cv2
import pickle
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from tqdm import tqdm
from sklearn.preprocessing import OneHotEncoder
from sklearn.metrics import confusion_matrix
from keras.models import Model, load_model
from keras.layers import Dense, Input, Conv2D, MaxPool2D, Flatten
from keras.preprocessing.image import ImageDataGenerator
from sklearn.model_selection import train_test_split
from keras.applications.vgg16 import VGG16
from keras.layers import GlobalAveragePooling2D
from keras.layers import BatchNormalization
from keras.layers import Dropout
from keras.layers import Dense
from keras.optimizers import Adam
from keras.callbacks import ReduceLROnPlateau
from keras.callbacks import ModelCheckpoint
from sklearn.metrics import ConfusionMatrixDisplay
import pandas as pd
```

```
[5]: def load_normal(norm_path):
    norm_files = np.array(os.listdir(norm_path))
    norm_labels = np.array(['normal']*len(norm_files))

    norm_images = []
    for image in tqdm(norm_files):
        image = cv2.imread(norm_path + image)
        image = cv2.resize(image, dsize=(200,200))
        image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
        norm_images.append(image)
```

```

norm_images = np.array(norm_images)
print(len(norm_images))

return norm_images, norm_labels
def load_pneumonia(pneu_path):
    pneu_files = np.array(os.listdir(pneu_path))
    pneu_labels = np.array([pneu_file.split('_')[1] for pneu_file in
↪pneu_files])

    pneu_images = []
    for image in tqdm(pneu_files):
        image = cv2.imread(pneu_path + image)
        image = cv2.resize(image, dsize=(200,200))

        image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
        pneu_images.append(image)

    pneu_images = np.array(pneu_images)
    X_train = np.append(norm_images, pneu_images, axis=0)
    y_train = np.append(norm_labels, pneu_labels)

    return pneu_images, pneu_labels

```

```

[6]: norm_images, norm_labels = load_normal('C:/Users/anask/OneDrive/Desktop/
↪Diplomski/3.sem/ASUB/chest_xray/chest_xray/train/NORMAL/')
pneu_images, pneu_labels = load_pneumonia('C:/Users/anask/OneDrive/Desktop/
↪Diplomski/3.sem/ASUB/chest_xray/chest_xray/train/PNEUMONIA/')

```

```
100%|      | 1341/1341 [00:28<00:00, 46.69it/s]
```

```
1341
```

```
100%|      | 3875/3875 [00:32<00:00, 119.76it/s]
```

```

[7]: X_train = np.append(norm_images, pneu_images, axis=0)
y_train = np.append(norm_labels, pneu_labels)

```

```

[8]: def plot_images(X, y):
    fig, axes = plt.subplots(ncols=7, nrows=2, figsize=(16, 4))

    indices = np.random.choice(len(X), 14)
    counter = 0

    for i in range(2):
        for j in range(7):
            axes[i,j].set_title(y[indices[counter]])
            axes[i,j].imshow(X[indices[counter]], cmap='gray')
            axes[i,j].get_xaxis().set_visible(False)

```

```
axes[i,j].get_yaxis().set_visible(False)
counter += 1
plt.show()
```

```
[9]: norm_images_test, norm_labels_test = load_normal('C:/Users/anask/OneDrive/
↳Desktop/Diplomski/3.sem/ASUB/chest_xray/chest_xray/test/NORMAL/')
pneu_images_test, pneu_labels_test = load_pneumonia('C:/Users/anask/OneDrive/
↳Desktop/Diplomski/3.sem/ASUB/chest_xray/chest_xray/test/PNEUMONIA/')

```

```
100%|      | 234/234 [00:04<00:00, 50.39it/s]
```

```
234
```

```
100%|      | 390/390 [00:03<00:00, 109.18it/s]
```


```
[10]: X_test = np.append(norm_images_test, pneu_images_test, axis=0)
y_test = np.append(norm_labels_test, pneu_labels_test)
```

```
[11]: with open('pneumonia_data.pickle', 'wb') as f:
        pickle.dump((X_train, X_test, y_train, y_test), f) # Use this to load
↳variables
with open('pneumonia_data.pickle', 'rb') as f:
        (X_train, X_test, y_train, y_test) = pickle.load(f)
```

```
[12]: np.unique(y_train, return_counts=True)
```

```
[12]: (array(['bacteria', 'normal', 'virus'], dtype='<U8'),
        array([2530, 1341, 1345], dtype=int64))
```


```
[13]: plot_images(X_train, y_train)
```



```
[14]: np.unique(y_train, return_counts=True)
```

```
[14]: (array(['bacteria', 'normal', 'virus'], dtype='<U8'),
        array([2530, 1341, 1345], dtype=int64))
```

```
[15]: plot_images(X_test, y_test)
```



```
[16]: y_train = y_train[:, np.newaxis]
      y_test = y_test[:, np.newaxis]
```

```
[17]: one_hot_encoder = OneHotEncoder(sparse=False, handle_unknown = 'error')
```

```
[18]: y_train_one_hot = one_hot_encoder.fit_transform(y_train)
      y_test_one_hot = one_hot_encoder.transform(y_test)
```

```
[19]: X_train = X_train.reshape(X_train.shape[0], X_train.shape[1], X_train.shape[2], 1)
      X_test = X_test.reshape(X_test.shape[0], X_test.shape[1], X_test.shape[2], 1)
```

```
[20]: datagen = ImageDataGenerator(
        rotation_range = 10,
        zoom_range = 0.1,
        width_shift_range = 0.1,
        height_shift_range = 0.1)
```

```
[21]: datagen.fit(X_train)
      train_gen = datagen.flow(X_train, y_train_one_hot, batch_size=32)
```

```
[22]: input1 = Input(shape=(X_train.shape[1], X_train.shape[2], 1))

cnn = Conv2D(16, (3, 3), activation='sigmoid', strides=(1, 1),
padding='same')(input1)
cnn = Conv2D(32, (3, 3), activation='sigmoid', strides=(1, 1),
padding='same')(cnn)
cnn = MaxPool2D((2, 2))(cnn)

cnn = Conv2D(16, (2, 2), activation='sigmoid', strides=(1, 1),
padding='same')(cnn)
cnn = Conv2D(32, (2, 2), activation='sigmoid', strides=(1, 1),
padding='same')(cnn)
```

```

cnn = MaxPool2D((2, 2))(cnn)

cnn = Flatten()(cnn)
cnn = Dense(100, activation='relu')(cnn)
cnn = Dense(50, activation='relu')(cnn)
output1 = Dense(3, activation='softmax')(cnn)

model = Model(inputs=input1, outputs=output1)

```

```
[23]: model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 200, 200, 1)]	0
conv2d (Conv2D)	(None, 200, 200, 16)	160
conv2d_1 (Conv2D)	(None, 200, 200, 32)	4640
max_pooling2d (MaxPooling2D)	(None, 100, 100, 32)	0
conv2d_2 (Conv2D)	(None, 100, 100, 16)	2064
conv2d_3 (Conv2D)	(None, 100, 100, 32)	2080
max_pooling2d_1 (MaxPooling2D)	(None, 50, 50, 32)	0
flatten (Flatten)	(None, 80000)	0
dense (Dense)	(None, 100)	8000100
dense_1 (Dense)	(None, 50)	5050
dense_2 (Dense)	(None, 3)	153
Total params: 8,014,247		
Trainable params: 8,014,247		
Non-trainable params: 0		

```
[24]: model.compile(loss='categorical_crossentropy', optimizer='adam',
    ↪ metrics=['acc'])
```

```
[25]: history2 = model.fit_generator(train_gen, epochs=30, validation_data=(X_test,
    ↪ y_test_one_hot))
```

C:\Users\anask\AppData\Local\Temp\ipykernel_11724\2115395957.py:1: UserWarning:
`Model.fit_generator` is deprecated and will be removed in a future version.
Please use `Model.fit`, which supports generators.

```
history2 = model.fit_generator(train_gen, epochs=30, validation_data=(X_test,
y_test_one_hot))
```

Epoch 1/30

163/163 [=====] - 132s 804ms/step - loss: 1.8028 - acc: 0.4057 - val_loss: 1.1575 - val_acc: 0.2372

Epoch 2/30

163/163 [=====] - 145s 888ms/step - loss: 1.0995 - acc: 0.4446 - val_loss: 1.1019 - val_acc: 0.3878

Epoch 3/30

163/163 [=====] - 126s 772ms/step - loss: 1.0729 - acc: 0.4755 - val_loss: 1.1144 - val_acc: 0.3878

Epoch 4/30

163/163 [=====] - 124s 758ms/step - loss: 1.0598 - acc: 0.4816 - val_loss: 1.1083 - val_acc: 0.3878

Epoch 5/30

163/163 [=====] - 126s 770ms/step - loss: 1.0560 - acc: 0.4850 - val_loss: 1.0877 - val_acc: 0.3878

Epoch 6/30

163/163 [=====] - 133s 816ms/step - loss: 1.0559 - acc: 0.4850 - val_loss: 1.1029 - val_acc: 0.3878

Epoch 7/30

163/163 [=====] - 134s 822ms/step - loss: 1.0557 - acc: 0.4829 - val_loss: 1.1291 - val_acc: 0.3878

Epoch 8/30

163/163 [=====] - 124s 759ms/step - loss: 1.0542 - acc: 0.4850 - val_loss: 1.1045 - val_acc: 0.3878

Epoch 9/30

163/163 [=====] - 132s 812ms/step - loss: 1.0525 - acc: 0.4850 - val_loss: 1.0999 - val_acc: 0.3878

Epoch 10/30

163/163 [=====] - 126s 775ms/step - loss: 1.0519 - acc: 0.4850 - val_loss: 1.1193 - val_acc: 0.3878

Epoch 11/30

163/163 [=====] - 122s 749ms/step - loss: 1.0533 - acc: 0.4850 - val_loss: 1.1477 - val_acc: 0.3878

Epoch 12/30

163/163 [=====] - 123s 752ms/step - loss: 1.0503 - acc: 0.4850 - val_loss: 1.1152 - val_acc: 0.3878

Epoch 13/30

163/163 [=====] - 123s 752ms/step - loss: 1.0031 - acc: 0.5077 - val_loss: 1.0457 - val_acc: 0.3878

Epoch 14/30
163/163 [=====] - 123s 751ms/step - loss: 0.9427 - acc:
0.5635 - val_loss: 1.0931 - val_acc: 0.5417

Epoch 15/30
163/163 [=====] - 135s 827ms/step - loss: 0.9330 - acc:
0.5763 - val_loss: 1.0006 - val_acc: 0.5593

Epoch 16/30
163/163 [=====] - 133s 818ms/step - loss: 0.9130 - acc:
0.5828 - val_loss: 0.9134 - val_acc: 0.5865

Epoch 17/30
163/163 [=====] - 128s 785ms/step - loss: 0.8188 - acc:
0.6281 - val_loss: 0.7308 - val_acc: 0.6603

Epoch 18/30
163/163 [=====] - 132s 811ms/step - loss: 0.7228 - acc:
0.6798 - val_loss: 0.7003 - val_acc: 0.7340

Epoch 19/30
163/163 [=====] - 123s 757ms/step - loss: 0.7095 - acc:
0.6814 - val_loss: 0.7852 - val_acc: 0.7115

Epoch 20/30
163/163 [=====] - 123s 752ms/step - loss: 0.6906 - acc:
0.6959 - val_loss: 0.7671 - val_acc: 0.7260

Epoch 21/30
163/163 [=====] - 122s 746ms/step - loss: 0.6670 - acc:
0.7120 - val_loss: 0.6754 - val_acc: 0.7404

Epoch 22/30
163/163 [=====] - 121s 743ms/step - loss: 0.6861 - acc:
0.7023 - val_loss: 0.8050 - val_acc: 0.6779

Epoch 23/30
163/163 [=====] - 124s 760ms/step - loss: 0.6637 - acc:
0.7118 - val_loss: 0.6050 - val_acc: 0.7596

Epoch 24/30
163/163 [=====] - 121s 743ms/step - loss: 0.6475 - acc:
0.7255 - val_loss: 0.7985 - val_acc: 0.7212

Epoch 25/30
163/163 [=====] - 121s 743ms/step - loss: 0.6572 - acc:
0.7174 - val_loss: 0.5942 - val_acc: 0.7853

Epoch 26/30
163/163 [=====] - 121s 739ms/step - loss: 0.6412 - acc:
0.7291 - val_loss: 0.6861 - val_acc: 0.7099

Epoch 27/30
163/163 [=====] - 121s 742ms/step - loss: 0.6286 - acc:
0.7301 - val_loss: 0.6335 - val_acc: 0.7660


Epoch 28/30
163/163 [=====] - 121s 743ms/step - loss: 0.6224 - acc:
0.7343 - val_loss: 0.6487 - val_acc: 0.7772


Epoch 29/30
163/163 [=====] - 121s 743ms/step - loss: 0.6256 - acc:
0.7339 - val_loss: 0.5449 - val_acc: 0.8141

Epoch 30/30

163/163 [=====] - 122s 749ms/step - loss: 0.6243 - acc: 0.7362 - val_loss: 0.6014 - val_acc: 0.7965

```
[26]: plt.figure(figsize=(8,6))
plt.title('Accuracy scores')
plt.plot(history2.history['acc'])
plt.plot(history2.history['val_acc'])
plt.legend(['acc', 'val_acc'])
plt.show()
plt.figure(figsize=(8,6))
plt.title('Loss value')
plt.plot(history2.history['loss'])
plt.plot(history2.history['val_loss'])
plt.legend(['loss', 'val_loss'])
plt.show()
```





```
[27]: X_train, X_test, y_train_one_hot, y_test_one_hot = train_test_split(
      X_train, y_train_one_hot, test_size=0.2, random_state = 42)
```

```
[28]: vgg16 = VGG16(weights='imagenet', include_top=False)
      hdf5_save = 'VGG16_Model.hdf5'
      annealer = ReduceLROnPlateau(
          monitor='val_accuracy', factor=0.70, patience=5,
          verbose=1, min_lr=1e-4)

      checkpoint = ModelCheckpoint(hdf5_save, verbose=1, save_best_only=True)

      datagen2 = ImageDataGenerator(rotation_range=360,
                                   width_shift_range=0.2,
                                   height_shift_range=0.2,
                                   zoom_range=0.2,
                                   horizontal_flip=True,
                                   vertical_flip=True)

      datagen2.fit(X_train)
      train_gen2 = datagen.flow(X_train, y_train_one_hot, batch_size=32)
```

```

input = Input(shape=(X_train.shape[1], X_train.shape[2], 1))
x = Conv2D(3, (3, 3), padding='same')(input)

x = vgg16(x)

x = GlobalAveragePooling2D()(x)
x = BatchNormalization()(x)
x = Dropout(0.5)(x)
x = Dense(200, activation='relu')(x)
x = BatchNormalization()(x)
x = Dropout(0.5)(x)

output = Dense(3, activation='softmax', name='root')(x)
model2 = Model(input, output)

optimizer = Adam(lr=0.003, beta_1=0.9, beta_2=0.999,
                  epsilon=0.1, decay=0.0)

model2.compile(loss='categorical_crossentropy',
               optimizer=optimizer, metrics=['accuracy'])

model2.summary()

history = model2.fit_generator(train_gen2, epochs=30, validation_data=(X_test,
↪y_test_one_hot))

```

Model: "model_1"

Layer (type)	Output Shape	Param #
input_3 (InputLayer)	[(None, 200, 200, 1)]	0
conv2d_4 (Conv2D)	(None, 200, 200, 3)	30
vgg16 (Functional)	(None, None, None, 512)	14714688
global_average_pooling2d (GlobalAveragePooling2D)	(None, 512)	0
batch_normalization (BatchNormalization)	(None, 512)	2048
dropout (Dropout)	(None, 512)	0
dense_3 (Dense)	(None, 200)	102600

batch_normalization_1 (Batch Normalization)	(None, 200)	800
dropout_1 (Dropout)	(None, 200)	0
root (Dense)	(None, 3)	603

=====

Total params: 14,820,769
Trainable params: 14,819,345
Non-trainable params: 1,424

Epoch 1/30

C:\FER\anaconda\lib\site-packages\keras\optimizers\optimizer_v2\adam.py:117:

UserWarning: The `lr` argument is deprecated, use `learning_rate` instead.

super().__init__(name, **kwargs)

C:\Users\anask\AppData\Local\Temp\ipykernel_11724\3426482288.py:43: UserWarning:

`Model.fit_generator` is deprecated and will be removed in a future version.

Please use `Model.fit`, which supports generators.

history = model2.fit_generator(train_gen2, epochs=30, validation_data=(X_test, y_test_one_hot))

131/131 [=====] - 1334s 10s/step - loss: 1.2980 - accuracy: 0.5117 - val_loss: 5.9136 - val_accuracy: 0.2883

Epoch 2/30

131/131 [=====] - 1393s 11s/step - loss: 1.0132 - accuracy: 0.5642 - val_loss: 2.3422 - val_accuracy: 0.4636

Epoch 3/30

131/131 [=====] - 1296s 10s/step - loss: 0.8370 - accuracy: 0.6237 - val_loss: 2.1336 - val_accuracy: 0.4636

Epoch 4/30

131/131 [=====] - 1326s 10s/step - loss: 0.7122 - accuracy: 0.6678 - val_loss: 1.6864 - val_accuracy: 0.4473

Epoch 5/30

131/131 [=====] - 1324s 10s/step - loss: 0.6785 - accuracy: 0.6918 - val_loss: 0.5777 - val_accuracy: 0.7107

Epoch 6/30

131/131 [=====] - 1276s 10s/step - loss: 0.6349 - accuracy: 0.7047 - val_loss: 0.5579 - val_accuracy: 0.7079

Epoch 7/30

131/131 [=====] - 1381s 11s/step - loss: 0.5992 - accuracy: 0.7270 - val_loss: 1.8965 - val_accuracy: 0.4636

Epoch 8/30

131/131 [=====] - 1366s 10s/step - loss: 0.5766 - accuracy: 0.7354 - val_loss: 1.3455 - val_accuracy: 0.5996


Epoch 9/30


131/131 [=====] - 1347s 10s/step - loss: 0.5717 -

accuracy: 0.7505 - val_loss: 2.1108 - val_accuracy: 0.4895
 Epoch 10/30
 131/131 [=====] - 1364s 10s/step - loss: 0.5360 -
 accuracy: 0.7649 - val_loss: 0.9281 - val_accuracy: 0.6226
 Epoch 11/30
 131/131 [=====] - 1390s 11s/step - loss: 0.4965 -
 accuracy: 0.7898 - val_loss: 0.9634 - val_accuracy: 0.7098
 Epoch 12/30
 131/131 [=====] - 1390s 11s/step - loss: 0.4961 -
 accuracy: 0.7855 - val_loss: 11.1811 - val_accuracy: 0.2749
 Epoch 13/30
 131/131 [=====] - 1400s 11s/step - loss: 0.4860 -
 accuracy: 0.7895 - val_loss: 0.9414 - val_accuracy: 0.6734
 Epoch 14/30
 131/131 [=====] - 1372s 11s/step - loss: 0.4592 -
 accuracy: 0.7975 - val_loss: 0.8226 - val_accuracy: 0.7165
 Epoch 15/30
 131/131 [=====] - 1253s 10s/step - loss: 0.4346 -
 accuracy: 0.8186 - val_loss: 1.2035 - val_accuracy: 0.7203
 Epoch 16/30
 131/131 [=====] - 1253s 10s/step - loss: 0.4327 -
 accuracy: 0.8102 - val_loss: 0.9534 - val_accuracy: 0.6159
 Epoch 17/30
 131/131 [=====] - 1249s 10s/step - loss: 0.4533 -
 accuracy: 0.8090 - val_loss: 0.7326 - val_accuracy: 0.7117
 Epoch 18/30
 131/131 [=====] - 1250s 10s/step - loss: 0.4241 -
 accuracy: 0.8150 - val_loss: 0.9361 - val_accuracy: 0.7538
 Epoch 19/30
 131/131 [=====] - 1252s 10s/step - loss: 0.4181 -
 accuracy: 0.8195 - val_loss: 2.4895 - val_accuracy: 0.5709
 Epoch 20/30
 131/131 [=====] - 1245s 10s/step - loss: 0.4036 -
 accuracy: 0.8305 - val_loss: 0.6826 - val_accuracy: 0.6973
 Epoch 21/30
 131/131 [=====] - 1240s 9s/step - loss: 0.4011 -
 accuracy: 0.8286 - val_loss: 0.4755 - val_accuracy: 0.8046
 Epoch 22/30
 131/131 [=====] - 1268s 10s/step - loss: 0.3715 -
 accuracy: 0.8495 - val_loss: 7.3258 - val_accuracy: 0.3573
 Epoch 23/30
 131/131 [=====] - 1238s 9s/step - loss: 0.3832 -
 accuracy: 0.8399 - val_loss: 0.9725 - val_accuracy: 0.7462
 Epoch 24/30
 131/131 [=====] - 1238s 9s/step - loss: 0.3775 -
 accuracy: 0.8437 - val_loss: 0.5535 - val_accuracy: 0.7787
 Epoch 25/30
 131/131 [=====] - 1235s 9s/step - loss: 0.3726 -

```
accuracy: 0.8411 - val_loss: 1.4503 - val_accuracy: 0.6715
Epoch 26/30
131/131 [=====] - 1236s 9s/step - loss: 0.3428 -
accuracy: 0.8535 - val_loss: 0.4583 - val_accuracy: 0.8257
Epoch 27/30
131/131 [=====] - 1235s 9s/step - loss: 0.3276 -
accuracy: 0.8598 - val_loss: 0.4417 - val_accuracy: 0.7883
Epoch 28/30
131/131 [=====] - 1236s 9s/step - loss: 0.3403 -
accuracy: 0.8514 - val_loss: 0.7734 - val_accuracy: 0.7443
Epoch 29/30
131/131 [=====] - 1231s 9s/step - loss: 0.3255 -
accuracy: 0.8684 - val_loss: 0.4532 - val_accuracy: 0.8084
Epoch 30/30
131/131 [=====] - 1232s 9s/step - loss: 0.3094 -
accuracy: 0.8682 - val_loss: 0.9094 - val_accuracy: 0.7222
```

```
[34]: plt.figure(figsize=(8,6))
plt.title('Accuracy scores')
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.legend(['accuracy', 'val_accuracy'])
plt.show()
plt.figure(figsize=(8,6))
plt.title('Loss value')
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.legend(['loss', 'val_loss'])
plt.show()
```





```
[32]: from sklearn.metrics import confusion_matrix
```

```
[36]: out = model2.predict(X_test)
```

```
33/33 [=====] - 77s 2s/step
```


```
[41]: out
      outp = np.where(out > .5, 1, 0)
      outp
```

```
[41]: array([[1, 0, 0],
            [0, 1, 0],
            [1, 0, 0],
            ...,
            [1, 0, 0],
            [0, 1, 0],
            [1, 0, 0]])
```

```
[47]: from sklearn.metrics import ConfusionMatrixDisplay
```

```
cm = ConfusionMatrixDisplay.from_predictions(y_true=y_test_one_hot.  
    ↳argmax(axis=1), y_pred=outp.argmax(axis=1), colorbar=False,↳  
    ↳display_labels=['bacteria', 'normal', 'virus'])  
cm.ax_.set(xlabel='Predicted', ylabel='True')
```

```
[47]: [Text(0.5, 0, 'Predicted'), Text(0, 0.5, 'True')]
```



```
[ ]:
```